

10) Implement and apply optimization methods for neural networks (AdaGrad, RMSProp, Adam) on any relevant dataset.

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader

def get_data():
    transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,),
(0.5,))])
    trainset = torchvision.datasets.MNIST(root='./data', train=True, download=True,
transform=transform)
    testset = torchvision.datasets.MNIST(root='./data', train=False, download=True,
transform=transform)
    trainloader = DataLoader(trainset, batch_size=64, shuffle=True)
    testloader = DataLoader(testset, batch_size=1000, shuffle=False)
    return trainloader, testloader

class SimpleNN(nn.Module):
    def __init__(self):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(28 * 28, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 10)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = x.view(-1, 28 * 28)
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.fc3(x)
        return x

def train_model(optimizer_type, trainloader, testloader):
    model = SimpleNN()
```

```

criterion = nn.CrossEntropyLoss()
optimizer = optimizer_type(model.parameters(), lr=0.01)

for epoch in range(5):
    model.train()
    for images, labels in trainloader:
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
    print(f"Epoch {epoch+1} completed")

model.eval()
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total
print(f"Accuracy with {optimizer_type.__name__}: {accuracy:.2f}%")

trainloader, testloader = get_data()
optimizers = [optim.Adagrad, optim.RMSprop, optim.Adam]

for opt in optimizers:
    train_model(opt, trainloader, testloader)

```

Output:

Epoch 1 completed

Epoch 2 completed

Epoch 3 completed

Epoch 4 completed

Epoch 5 completed

Accuracy with Adagrad: 94.78%

Epoch 1 completed

Epoch 2 completed

Epoch 3 completed

Epoch 4 completed

Epoch 5 completed

Accuracy with RMSprop: 93.17%

Epoch 1 completed

Epoch 2 completed

Epoch 3 completed

Epoch 4 completed

Epoch 5 completed

Accuracy with Adam: 94.48%

- **Data preprocessing**
- **Model definition (simple feedforward network)**
- **Training using different optimizers**
- **Performance comparison**

Apply, train and visualize Different deep CNN architectures like LeNet, AlexNet, VGG, PlacesNet, on MNIST datasets.

This program uses LeNet, AlexNet on MNIST dataset

```
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import numpy as np

# Load MNIST dataset
def get_data():
    transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,),
(0.5,))])
    trainset = torchvision.datasets.MNIST(root='./data', train=True, download=True,
transform=transform)
    testset = torchvision.datasets.MNIST(root='./data', train=False, download=True,
transform=transform)
    trainloader = DataLoader(trainset, batch_size=64, shuffle=True)
    testloader = DataLoader(testset, batch_size=1000, shuffle=False)
    return trainloader, testloader

# Define CNN Architectures
class LeNet(nn.Module):
    def __init__(self):
        super(LeNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 6, kernel_size=5)
        self.conv2 = nn.Conv2d(6, 16, kernel_size=5)
        self.fc1 = nn.Linear(16*4*4, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
        self.relu = nn.ReLU()
        self.pool = nn.MaxPool2d(2, 2)

    def forward(self, x):
        x = self.pool(self.relu(self.conv1(x)))
        x = self.pool(self.relu(self.conv2(x)))
```

```

x = x.view(-1, 16*4*4)
x = self.relu(self.fc1(x))
x = self.relu(self.fc2(x))
x = self.fc3(x)
return x

```

```

class AlexNet(nn.Module):
    def __init__(self):
        super(AlexNet, self).__init__()
        self.features = nn.Sequential(
            nn.Conv2d(1, 64, kernel_size=3, stride=1, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(64, 192, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Conv2d(192, 384, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.Conv2d(384, 256, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2),
        )
        self.classifier = nn.Sequential(
            nn.Linear(256 * 3 * 3, 4096),
            nn.ReLU(),
            nn.Linear(4096, 4096),
            nn.ReLU(),
            nn.Linear(4096, 10),
        )

    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x

```

Train and evaluate model

```

def train_model(model, trainloader, testloader, optimizer_type):
    model = model()
    criterion = nn.CrossEntropyLoss()
    optimizer = optimizer_type(model.parameters(), lr=0.01)

```

```

for epoch in range(5):
    model.train()
    for images, labels in trainloader:
        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
    print(f"Epoch {epoch+1} completed")

model.eval()
correct = 0
total = 0
with torch.no_grad():
    for images, labels in testloader:
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

accuracy = 100 * correct / total
print(f"Accuracy with {model.__class__.__name__}: {accuracy:.2f}%")
return model

```

Visualization Function

```

def visualize_filters(model):
    with torch.no_grad():
        for name, param in model.named_parameters():
            if 'conv' in name and param.requires_grad:
                filters = param.cpu().numpy()
                fig, axes = plt.subplots(1, min(6, filters.shape[0]))
                for i, ax in enumerate(axes):
                    ax.imshow(filters[i, 0], cmap='gray')
                    ax.axis('off')
                plt.show()
                break

```

Load data

```

trainloader, testloader = get_data()

```

Train models

```

models = [LeNet, AlexNet]
optimizers = [optim.Adam]
for model in models:

```

for opt in optimizers:

```
    trained_model = train_model(model, trainloader, testloader, opt)
    visualize_filters(trained_model)
```

```
100%|██████████| 9.91M/9.91M [00:00<00:00, 55.8MB/s]
100%|██████████| 28.9k/28.9k [00:00<00:00, 1.70MB/s]
100%|██████████| 1.65M/1.65M [00:00<00:00, 14.0MB/s]
100%|██████████| 4.54k/4.54k [00:00<00:00, 6.21MB/s]
```

Epoch 1 completed


Epoch 2 completed

Epoch 3 completed

Epoch 4 completed

Epoch 5 completed

Accuracy with LeNet: 98.17%



Epoch 1 completed

Epoch 2 completed

Epoch 3 completed

Epoch 4 completed

Epoch 5 completed

Accuracy with AlexNet: 10.09%

This program uses VGG, PlacesNet on MNIST dataset

```
import tensorflow as tf
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Input
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
```

```

import cv2

# Load MNIST Dataset
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

# Preprocess Data
x_train = np.expand_dims(x_train, axis=-1)
x_test = np.expand_dims(x_test, axis=-1)

# Convert grayscale to 3 channels for pretrained models
x_train = np.repeat(x_train, 3, axis=-1)
x_test = np.repeat(x_test, 3, axis=-1)

# Resize images to 32x32 for VGG-16 compatibility
x_train_resized = np.array([cv2.resize(img, (32, 32)) for img in x_train])
x_test_resized = np.array([cv2.resize(img, (32, 32)) for img in x_test])

# Normalize data
x_train_resized, x_test_resized = x_train_resized / 255.0, x_test_resized / 255.0
x_train, x_test = x_train / 255.0, x_test / 255.0

# Convert labels to categorical
y_train_cat = to_categorical(y_train, 10)
y_test_cat = to_categorical(y_test, 10)

# Define VGG-16 Model
base_model_vgg = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
for layer in base_model_vgg.layers:
    layer.trainable = False

x = Flatten()(base_model_vgg.output)
x = Dense(256, activation='relu')(x)
x = Dense(10, activation='softmax')(x)

vgg_model = Model(inputs=base_model_vgg.input, outputs=x)
vgg_model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])

# Train VGG-16
vgg_model.fit(x_train_resized, y_train_cat, epochs=5, batch_size=64,
validation_data=(x_test_resized, y_test_cat))

```



```

# Define PlacesNet-like CNN
input_layer = Input(shape=(28, 28, 3))
x = Conv2D(64, (3,3), activation='relu', padding='same')(input_layer)
x = MaxPooling2D((2,2))(x)
x = Conv2D(128, (3,3), activation='relu', padding='same')(x)
x = MaxPooling2D((2,2))(x)
x = Flatten()(x)
x = Dense(256, activation='relu')(x)
x = Dense(10, activation='softmax')(x)

placesnet_model = Model(inputs=input_layer, outputs=x)
placesnet_model.compile(optimizer=Adam(), loss='categorical_crossentropy',
metrics=['accuracy'])

# Train PlacesNet
placesnet_model.fit(x_train, y_train_cat, epochs=5, batch_size=64, validation_data=(x_test,
y_test_cat))

# Evaluate and Visualize Results
def plot_confusion_matrix(model, x_test, y_test, title):
    y_pred = np.argmax(model.predict(x_test), axis=1)
    cm = confusion_matrix(y_test, y_pred)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.title(f'Confusion Matrix: {title}')
    plt.show()

# Confusion Matrices
plot_confusion_matrix(vgg_model, x_test_resized, y_test, "VGG-16")
plot_confusion_matrix(placesnet_model, x_test, y_test, "PlacesNet")

# Display Feature Maps
def visualize_feature_maps(model, x_sample):
    layer_outputs = [layer.output for layer in model.layers if isinstance(layer, Conv2D)]
    activation_model = Model(inputs=model.input, outputs=layer_outputs)
    activations = activation_model.predict(np.expand_dims(x_sample, axis=0))
    for i, activation in enumerate(activations[:3]): # Show first 3 layers

```

```
plt.figure(figsize=(10, 5))
for j in range(min(activation.shape[-1], 6)): # Show first 6 filters
    plt.subplot(1, 6, j+1)
    plt.imshow(activation[0, :, :, j], cmap='viridis')
    plt.axis('off')
plt.show()


# Show feature maps of first test image
visualize_feature_maps(vgg_model, x_test_resized[0])
visualize_feature_maps(placesnet_model, x_test[0])
```

Output:

938/938

 204s
205ms/step - accuracy: 0.9961 - loss: 0.0109 - val_accuracy: 0.9891 -
val_loss: 0.0339

Confusion Matrix: VGG-16



Confusion Matrix: PlacesNet

